

```
In [2]: import pandas as pd
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score
import seaborn as sns

In [3]: df = pd.read_csv("CarPrice_Assignment.csv")

In [4]: x=df[['engineize','horsepower','citympg','highwaympg']]
y=df['price']

In [5]: X_train, X_test, y_train, y_test = train_test_split(x, y, test_size=0.2, random_state=42)

In [6]: df = df.dropna()

In [7]: print(df.head())

   car_ID  symboling      CarName  fueltpe aspiration  doornumber  \
0         1           3    alfa-romero giulia         gas         std         two
1         2           3    alfa-romero stelvio         gas         std         two
2         3           1  alfa-romero Quadrifoglio         gas         std         two
3         4           2      audi 100 ls         gas         std         four
4         5           2      audi 100ls         gas         std         four

   carbody drivewheel  enginelocation  wheelbase  ...  enginesize  \
0  convertible         rwd         front      88.6  ...        130
1  convertible         rwd         front      88.6  ...        130
2   hatchback         rwd         front      94.5  ...        152
3     sedan         fwd         front      99.8  ...        109
4     sedan         4wd         front      99.4  ...        136

   fuelsystem  boreratio  stroke  compressionratio  horsepower  peakrpm  citympg  \
0      mpfi         3.47    2.68             9.0         111    5000      21
1      mpfi         3.47    2.68             9.0         111    5000      21
2      mpfi         2.68    3.47             9.0         154    5000      19
3      mpfi         3.19    3.40            10.0         102    5500      24
4      mpfi         3.19    3.40             8.0         115    5500      18

   highwaympg  price
0             27  13495.0
1             27  16500.0
2             26  16500.0
3             30  13950.0
4             22  17450.0

[5 rows x 26 columns]
```

```
In [8]: df
```

	car_ID	symboling	CarName	fueltpe	aspiration	doornumber	carbody	drivewheel	enginelocation	wheelbase	...	engineize	fuelsystem	boreratio	stroke	compressionratio	horsepower	peakrpm	citympg	highwaym
0	1	3	alfa-romero giulia	gas	std	two	convertible	rwd	front	88.6	...	130	mpfi	3.47	2.68	9.0	111	5000	21	
1	2	3	alfa-romero stelvio	gas	std	two	convertible	rwd	front	88.6	...	130	mpfi	3.47	2.68	9.0	111	5000	21	
2	3	1	alfa-romero Quadrifoglio	gas	std	two	hatchback	rwd	front	94.5	...	152	mpfi	2.68	3.47	9.0	154	5000	19	
3	4	2	audi 100 ls	gas	std	four	sedan	fwd	front	99.8	...	109	mpfi	3.19	3.40	10.0	102	5500	24	
4	5	2	audi 100ls	gas	std	four	sedan	4wd	front	99.4	...	136	mpfi	3.19	3.40	8.0	115	5500	18	
...
200	201	-1	volvo 145e (sw)	gas	std	four	sedan	rwd	front	109.1	...	141	mpfi	3.78	3.15	9.5	114	5400	23	
201	202	-1	volvo 144ea	gas	turbo	four	sedan	rwd	front	109.1	...	141	mpfi	3.78	3.15	8.7	160	5300	19	
202	203	-1	volvo 244dl	gas	std	four	sedan	rwd	front	109.1	...	173	mpfi	3.58	2.87	8.8	134	5500	18	
203	204	-1	volvo 246	diesel	turbo	four	sedan	rwd	front	109.1	...	145	idi	3.01	3.40	23.0	106	4800	26	
204	205	-1	volvo 264gl	gas	turbo	four	sedan	rwd	front	109.1	...	141	mpfi	3.78	3.15	9.5	114	5400	19	

205 rows x 26 columns

```
In [9]: SS = StandardScaler()

x_train = SS.fit_transform(X_train)
x_test = SS.fit_transform(X_test)

In [10]: model = LinearRegression()
model.fit(x_train, y_train)

Out[10]: ▼ LinearRegression
LinearRegression()

In [11]: y_pred=model.predict(x_test)

In [12]: print("-"*50)
print("MODEL COEFFICIENTS:")
for feature, coef in zip(x.columns, model.coef_):
    print(f"({feature}>12): (coef:>10.2f)")
print(f"('Intercept':>12): (model.intercept_:>10.2f)")

-----
MODEL COEFFICIENTS:
(feature:>12): (coef:>10.2f)
(feature:>12): (coef:>10.2f)
(feature:>12): (coef:>10.2f)
(feature:>12): (coef:>10.2f)
('Intercept':>12): (model.intercept_:>10.2f)
```

```
In [13]: import numpy as np
from sklearn.metrics import mean_squared_error, r2_score

print("\nMODEL PERFORMANCE:")

print(f"{'MSE':>12}: (mean_squared_error(y_test, y_pred):>10.2f)")
print(f"{'RMSE':>12}: {np.sqrt(mean_squared_error(y_test, y_pred)):>10.2f}")
print(f"{'R-squared':>12}: (r2_score(y_test, y_pred):>10.2f)")


print("-" * 58)

MODEL PERFORMANCE:
      MSE: 17133687.82
     RMSE: 4139.29
  R-squared: 0.78
=====
```

```
In [14]: mse=mean_squared_error(y_test,y_pred)
print('Mean Squared Error=',mse)
rmse=np.sqrt(mse)
print("Root Mean Squared Error=",rmse)
r2score=r2_score(y_test,y_pred)
print('r2 score=',r2score)

Mean Squared Error= 17133687.822608914
Root Mean Squared Error= 4139.285907328572
r2 score= 0.7829640847819556
```


```
In [15]: plt.figure(figsize=(10, 5))
plt.scatter(y_test, y_pred, alpha=0.6)
plt.plot([y.min(), y.max()], [y.min(), y.max()], 'r--')
plt.title("Linearity Check: Actual vs Predicted Prices")
plt.xlabel("Actual Price ($)")
plt.ylabel("Predicted Price ($)")
plt.grid(True)
plt.show()
```



```
In [20]: #2. Independence (Durbin-Watson)
import statsmodels.api as sm
residuals = y_test - y_pred
dw_test = sm.stats.durbin_watson(residuals)
print(f"\nDurbin-Watson Statistic: {dw_test:.2f}",
      "\n(Values close to 2 indicate no autocorrelation)")

# 3. Homoscedasticity
plt.figure(figsize=(10, 5))
sns.residplot(x=y_pred, y=residuals, lowess=True, line_kws={'color': 'red'})
plt.title("Homoscedasticity Check: Residuals vs Predicted")
plt.xlabel("Predicted Price ($)")
plt.ylabel("Residuals ($)")
plt.grid(True)
plt.show()
```

Durbin-Watson Statistic: 2.34
(Values close to 2 indicate no autocorrelation)



```
In [21]: #4. Normality of residuals
fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(12,5))
sns.histplot(residuals, kde=True, ax=ax1)
ax1.set_title("Residuals Distribution")
sm.qqplot(residuals, line='45', fit=True, ax=ax2)
ax2.set_title("Q-Q Plot")
plt.tight_layout()
plt.show()
```

